

Advanced Cloud Computing

MapReduce Algorithm Design

Wei Wang
CSE@HKUST
Spring 2025



THE DEPARTMENT OF

COMPUTER SCIENCE & ENGINEERING

計算機科學及工程學系

Limited control

All algorithms must be expressed in **m**, **r**, **c**, **p**

You don't know

- ▶ where mappers and reducers run
- ▶ when a mapper or reducer begins or finishes
- ▶ which input a particular mapper is processing
- ▶ which intermediate key a particular reducer is processing

But still we can control

Cleverly-constructed data structures

- ▶ bring partial results together

Sort order of intermediate keys

- ▶ control order in which reducers process keys

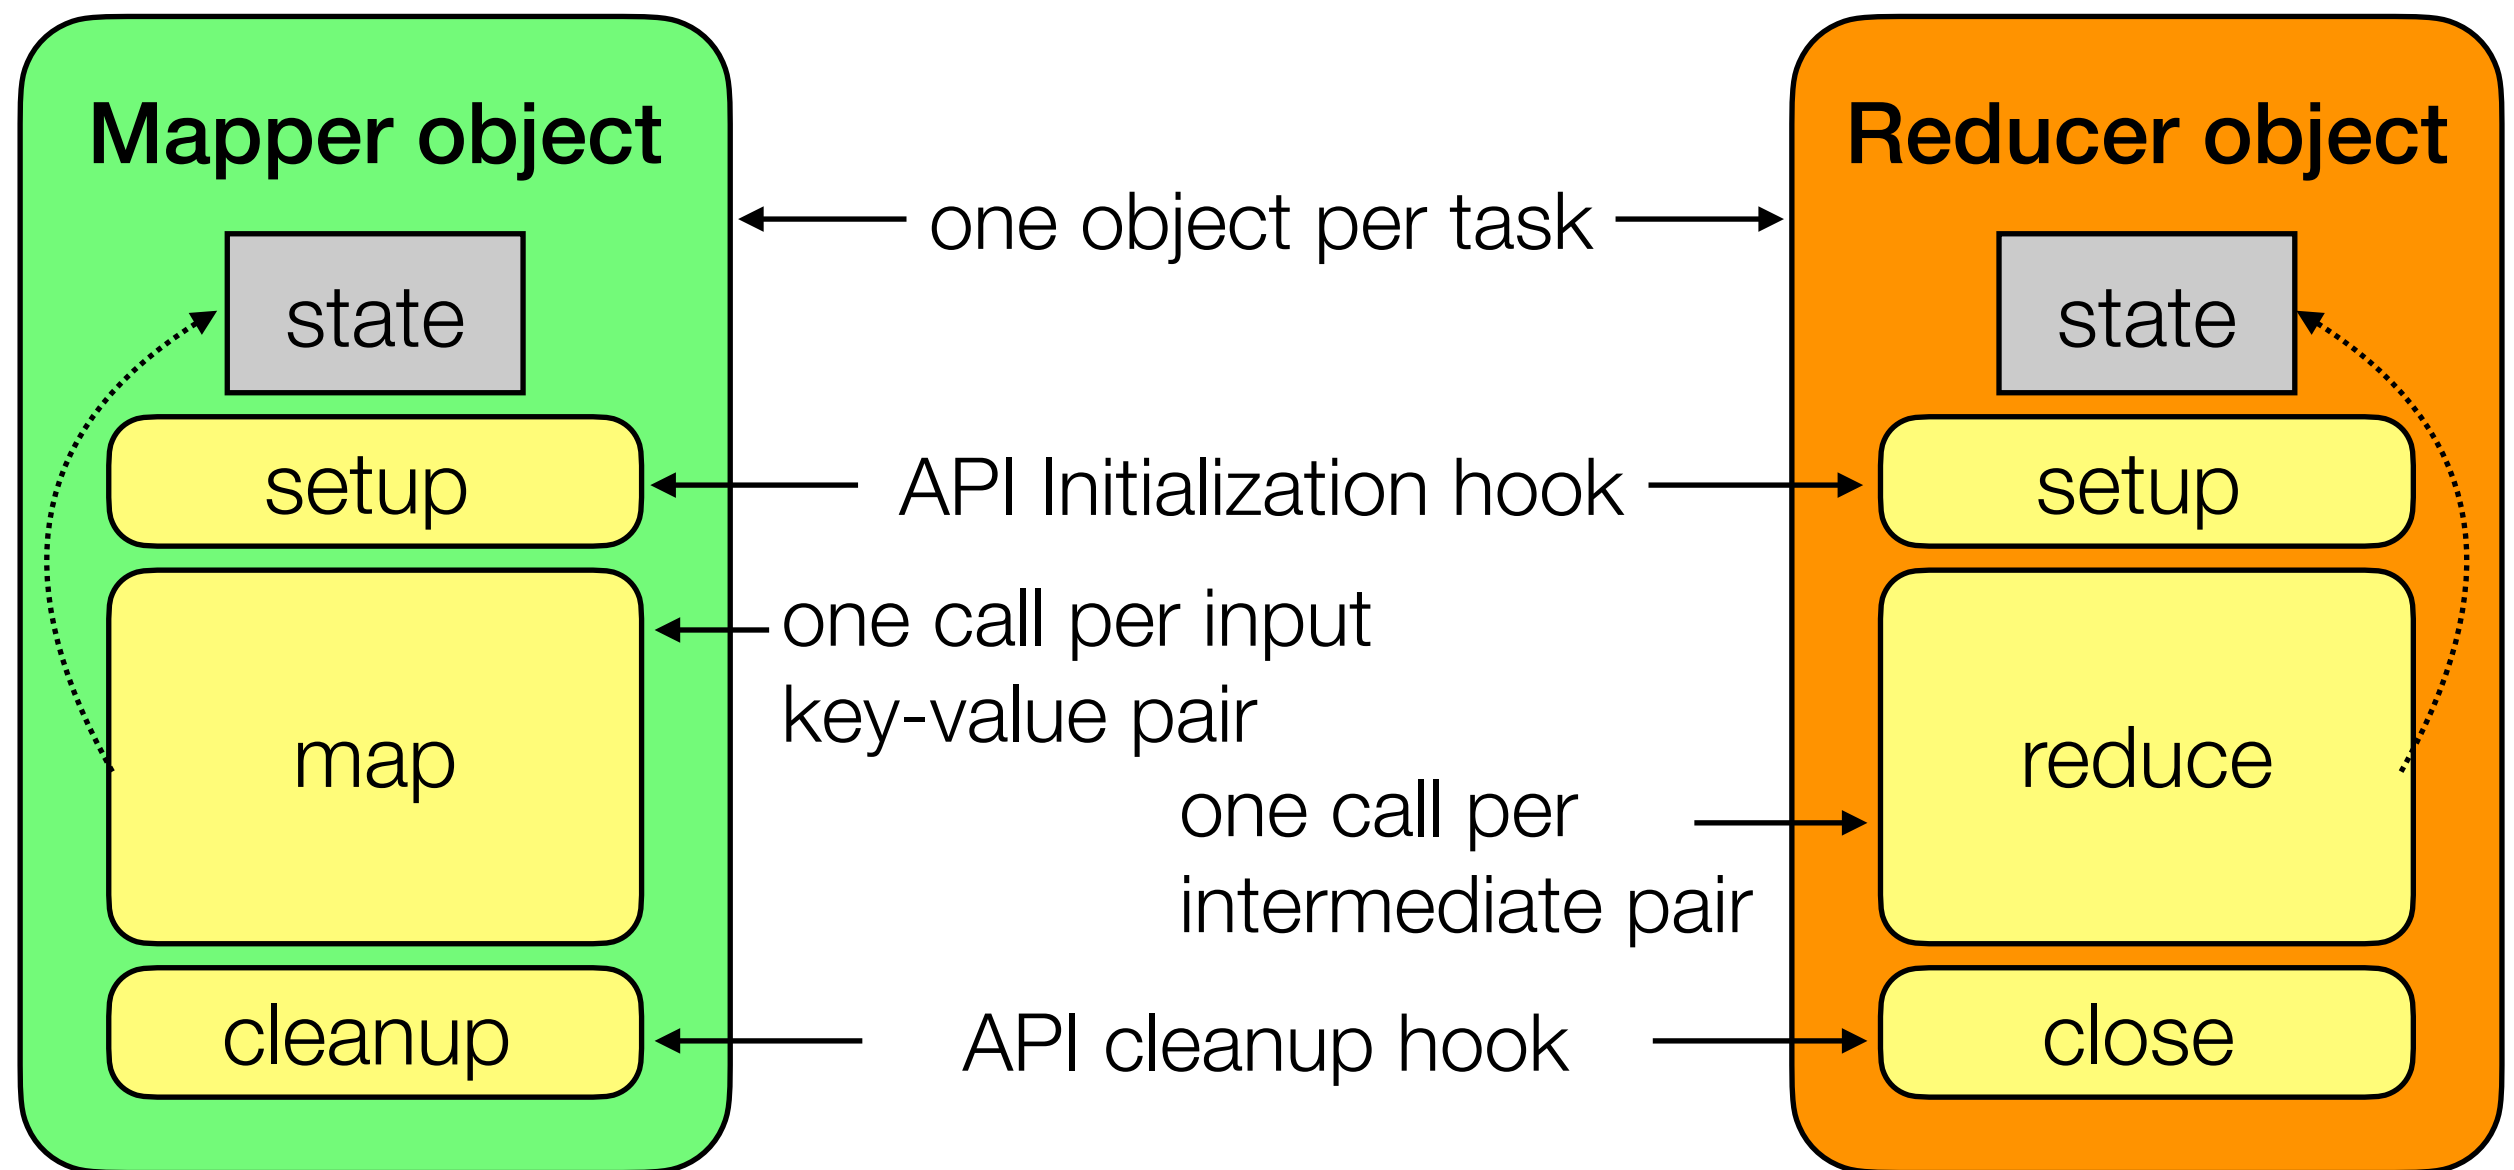
Partitioner

- ▶ control which reducer processes which keys

Preserving state in mappers and reducers

- ▶ capture dependencies across multiple keys and values

Preserving state



Scalable Hadoop Algorithms

Avoid object creation

- ▶ inherently costly operation
- ▶ garbage collection

Avoid buffering

- ▶ limited heap size
- ▶ works for small datasets, but won't scale!

Importance of local aggregation

Ideal scaling characteristics

- ▶ twice the data, twice the running time
- ▶ twice the resources, half the running time

Why can't we achieve this?

- ▶ synchronization requires communication
- ▶ communication kills performance

Thus... avoid communication, as much as possible!

- ▶ reduce intermediate data via local aggregation
- ▶ combiners can help

WordCount: Baseline

```
1: class MAPPER
2:   method MAP(docid  $a$ , doc  $d$ )
3:     for all term  $t \in \text{doc } d$  do
4:       EMIT(term  $t$ , count 1)

1: class REDUCER
2:   method REDUCE(term  $t$ , counts  $[c_1, c_2, \dots]$ )
3:      $sum \leftarrow 0$ 
4:     for all count  $c \in \text{counts } [c_1, c_2, \dots]$  do
5:        $sum \leftarrow sum + c$ 
6:     EMIT(term  $t$ , count  $sum$ )
```

What's the impact of combiners?

WordCount: Version 1

$H\{t\}$: a hash table

```
1: class MAPPER
2:   method MAP(docid  $a$ , doc  $d$ )
3:      $H \leftarrow$  new ASSOCIATIVEARRAY
4:     for all term  $t \in$  doc  $d$  do
5:        $H\{t\} \leftarrow H\{t\} + 1$             $\triangleright$  Tally counts for entire document
6:     for all term  $t \in H$  do
7:       EMIT(term  $t$ , count  $H\{t\}$ )
```

Do combiners still help?

WordCount: Version 2

$H\{t\}$: a hash table

Key idea: preserve state across
input key-value pairs!

```
1: class MAPPER
2:   method INITIALIZE
3:      $H \leftarrow \text{new ASSOCIATIVEARRAY}$ 
4:   method MAP(docid  $a$ , doc  $d$ )
5:     for all term  $t \in \text{doc } d$  do
6:        $H\{t\} \leftarrow H\{t\} + 1$            ▷ Tally counts across documents
7:   method CLOSE
8:     for all term  $t \in H$  do
9:       EMIT(term  $t$ , count  $H\{t\}$ )
```

Do combiners still help?

Design pattern for local aggregation

In-mapper combining

- ▶ fold the functionality of the combiner into the mapper by preserving state across multiple map calls

Advantages

- ▶ speed
- ▶ why is this faster than actual combiners?

Disadvantages

- ▶ explicit memory management required
- ▶ Order matters! May lead to order-dependent bugs!

Combiner design

Combiners and reducers share same method signature

- ▶ sometimes, reducers can serve combiners
- ▶ often, not...

Combiners are **optional** optimizations

- ▶ should not affect algorithm correctness
- ▶ may be run 0, 1, or multiple times (indefinite)

Example: find the mean of integers associated with the same key

Computing the mean: Version 1

```
1: class MAPPER
2:   method MAP(string  $t$ , integer  $r$ )
3:     EMIT(string  $t$ , integer  $r$ )

1: class REDUCER
2:   method REDUCE(string  $t$ , integers  $[r_1, r_2, \dots]$ )
3:      $sum \leftarrow 0$ 
4:      $cnt \leftarrow 0$ 
5:     for all integer  $r \in$  integers  $[r_1, r_2, \dots]$  do
6:        $sum \leftarrow sum + r$ 
7:        $cnt \leftarrow cnt + 1$ 
8:      $r_{avg} \leftarrow sum / cnt$ 
9:     EMIT(string  $t$ , integer  $r_{avg}$ )
```

Why can't we use reducer as combiner?

Computing the mean: Version 2

```
1: class MAPPER
2:   method MAP(string  $t$ , integer  $r$ )
3:     EMIT(string  $t$ , integer  $r$ )

1: class COMBINER
2:   method COMBINE(string  $t$ , integers  $[r_1, r_2, \dots]$ )
3:      $sum \leftarrow 0$ 
4:      $cnt \leftarrow 0$ 
5:     for all integer  $r \in$  integers  $[r_1, r_2, \dots]$  do
6:        $sum \leftarrow sum + r$ 
7:        $cnt \leftarrow cnt + 1$ 
8:     EMIT(string  $t$ , pair ( $sum, cnt$ )) ▷ Separate sum and count

1: class REDUCER
2:   method REDUCE(string  $t$ , pairs  $[(s_1, c_1), (s_2, c_2) \dots]$ )
3:      $sum \leftarrow 0$ 
4:      $cnt \leftarrow 0$ 
5:     for all pair  $(s, c) \in$  pairs  $[(s_1, c_1), (s_2, c_2) \dots]$  do
6:        $sum \leftarrow sum + s$ 
7:        $cnt \leftarrow cnt + c$ 
8:      $r_{avg} \leftarrow sum / cnt$ 
9:     EMIT(string  $t$ , integer  $r_{avg}$ )
```

Why doesn't this work?

Computing the mean: Version 3

```
1: class MAPPER
2:   method MAP(string  $t$ , integer  $r$ )
3:     EMIT(string  $t$ , pair ( $r$ , 1))
```

Fixed?

```
1: class COMBINER
2:   method COMBINE(string  $t$ , pairs  $[(s_1, c_1), (s_2, c_2) \dots]$ )
3:      $sum \leftarrow 0$ 
4:      $cnt \leftarrow 0$ 
5:     for all pair  $(s, c) \in$  pairs  $[(s_1, c_1), (s_2, c_2) \dots]$  do
6:        $sum \leftarrow sum + s$ 
7:        $cnt \leftarrow cnt + c$ 
8:     EMIT(string  $t$ , pair ( $sum$ ,  $cnt$ ))
```

```
1: class REDUCER
2:   method REDUCE(string  $t$ , pairs  $[(s_1, c_1), (s_2, c_2) \dots]$ )
3:      $sum \leftarrow 0$ 
4:      $cnt \leftarrow 0$ 
5:     for all pair  $(s, c) \in$  pairs  $[(s_1, c_1), (s_2, c_2) \dots]$  do
6:        $sum \leftarrow sum + s$ 
7:        $cnt \leftarrow cnt + c$ 
8:      $r_{avg} \leftarrow sum / cnt$ 
9:     EMIT(string  $t$ , integer  $r_{avg}$ )
```

Computing the mean: Version 4

```
1: class MAPPER
2:   method INITIALIZE
3:      $S \leftarrow \text{new ASSOCIATIVEARRAY}$ 
4:      $C \leftarrow \text{new ASSOCIATIVEARRAY}$ 
5:   method MAP(string  $t$ , integer  $r$ )
6:      $S\{t\} \leftarrow S\{t\} + r$ 
7:      $C\{t\} \leftarrow C\{t\} + 1$ 
8:   method CLOSE
9:     for all term  $t \in S$  do
10:      EMIT(term  $t$ , pair ( $S\{t\}$ ,  $C\{t\}$ ))
```

Are combiners still needed?

Algorithm design: a running example

“Pairs” approach

Each mapper takes a sentence:

- ▶ generate all co-occurring term pairs
- ▶ for all pairs, emit $(a,b) \rightarrow \text{count}$

Reducers sum up counts associated with these pairs

Use combiners to minimize shuffling traffics!

“Pairs”: pseudo-code

```
1: class MAPPER
2:   method MAP(docid  $a$ , doc  $d$ )
3:     for all term  $w \in \text{doc } d$  do
4:       for all term  $u \in \text{NEIGHBORS}(w)$  do
5:         EMIT(pair  $(w, u)$ , count 1)           ▷ Emit count for each
co-occurrence
1: class REDUCER
2:   method REDUCE(pair  $p$ , counts  $[c_1, c_2, \dots]$ )
3:      $s \leftarrow 0$ 
4:     for all count  $c \in \text{counts } [c_1, c_2, \dots]$  do
5:        $s \leftarrow s + c$                        ▷ Sum co-occurrence counts
6:     EMIT(pair  $p$ , count  $s$ )
```

“Pairs” analysis

Advantages

- ▶ easy to implement, easy to understand

Disadvantages

- ▶ lots of pairs to sort and shuffle around
 - ▶ What's the upper bound?
- ▶ not many opportunities for combiners to work

“Stripes” approach

Group together pairs into an associative array

$(a, b) \rightarrow 1$

$(a, c) \rightarrow 2$

$(a, d) \rightarrow 5$

$(a, e) \rightarrow 3$

$(a, f) \rightarrow 2$

$a \rightarrow \{b: 1, c: 2, d: 5, e: 3, f: 2\}$

Each mapper takes a sentence

- ▶ generate all co-occurring term pairs
- ▶ for each term, emit $a \rightarrow \{ b: \text{count}_b, c: \text{count}_c, d: \text{count}_d \dots \}$

“Stripes” approach

Reducers perform

- ▶ **element-wise sum** of associative arrays

$$\begin{array}{r} a \rightarrow \{ b: 1, \quad d: 5, e: 3 \} \\ + \quad a \rightarrow \{ b: 1, c: 2, d: 2, \quad f: 2 \} \\ \hline a \rightarrow \{ b: 2, c: 2, d: 7, e: 3, f: 2 \} \end{array}$$

Key idea: cleverly-constructed data structure brings together partial results

“Stripes”: pseudo-code

```
1: class MAPPER
2:   method MAP(docid  $a$ , doc  $d$ )
3:     for all term  $w \in \text{doc } d$  do
4:        $H \leftarrow \text{new ASSOCIATIVEARRAY}$ 
5:       for all term  $u \in \text{NEIGHBORS}(w)$  do
6:          $H\{u\} \leftarrow H\{u\} + 1$   $\triangleright$  Tally words co-occurring with  $w$ 
7:       EMIT(Term  $w$ , Stripe  $H$ )

1: class REDUCER
2:   method REDUCE(term  $w$ , stripes  $[H_1, H_2, H_3, \dots]$ )
3:      $H_f \leftarrow \text{new ASSOCIATIVEARRAY}$ 
4:     for all stripe  $H \in \text{stripes } [H_1, H_2, H_3, \dots]$  do
5:       SUM( $H_f, H$ )  $\triangleright$  Element-wise sum
6:     EMIT(term  $w$ , stripe  $H_f$ )
```

“Stripes” analysis

Advantages

- ▶ far less sorting and shuffling of key-value pairs
- ▶ can make better use of combiners

Disadvantages

- ▶ more difficult to implement
- ▶ underlying object more heavyweight
- ▶ fundamental limitation in terms of size of event space

Relative frequencies

How do we estimate relative frequencies from counts?

$$f(B|A) = \frac{N(A, B)}{N(A)} = \frac{N(A, B)}{\sum_{B'} N(A, B')}$$

Why do we want to do this?

How do we do this with MapReduce?

$f(B|A)$: “Stripes”

$a \rightarrow \{b: 1, c: 2, d: 5, e: 3, f: 2\}$

Easy!

- ▶ one pass to compute $(a, *)$
- ▶ another pass to directly compute $f(B|A)$

$f(B|A)$: “Pairs”

$(a, b) \rightarrow 1$

$(a, c) \rightarrow 2$

$(a, d) \rightarrow 5$

$(a, e) \rightarrow 3$

...

What're the issues?


- ▶ computing relative frequencies requires **marginal counts**
- ▶ but the marginal cannot be computed until you see all counts

Buffering is a bad idea! (why?)

Solution: what if we could get the marginal count to arrive at the reducer first?

$f(B|A)$: “Pairs”

$(a, *) \rightarrow 20$ Reducer holds this value in memory




Emit extra $(a, *)$ for every “b”, “c”, “d”, “e”... in mapper

Make sure all a’s get sent to the same reducer (partitioner)

Make sure $(a, *)$ comes first (define sort order)

Hold state in reducer across different key-value pairs

From a reducer's perspective



key	values	
(dog, *)	[6327, 8514, ...]	compute marginal: $\sum_{w'} N(\text{dog}, w') = 42908$
(dog, aardvark)	[2, 1]	$f(\text{aardvark} \text{dog}) = 3/42908$
(dog, aardwolf)	[1]	$f(\text{aardwolf} \text{dog}) = 1/42908$
...		
(dog, zebra)	[2, 1, 1, 1]	$f(\text{zebra} \text{dog}) = 5/42908$
(doge, *)	[682, ...]	compute marginal: $\sum_{w'} N(\text{doge}, w') = 1267$

What to emit in mapper?

Emit extra (a, *) for every “b”, “c”, “d”, “e”... in mapper

```
1: class MAPPER
2:     method MAP(docid  $a$ , doc  $d$ )
3:         for all term  $w \in \text{doc } d$  do
4:             for all term  $u \in \text{NEIGHBORS}(w)$  do
5:                 EMIT(pair ( $w, u$ ), count 1)
6:                 Emit(pair ( $w, ""$ ), count 1)
```

How to make sure that pair ($w, ""$) is sorted in order before pair (w, u)?

Write your own data types

`PairOfStrings` implements `WritableComparable` interface

Must implement

- ▶ `write` for serialization
- ▶ `readFields` for deserialization
- ▶ `compareTo` to define sort order

Sample code available in Assignment-3

Partitioner

Make sure all a's get sent to the same reducer (use partitioner)

```
private static class MyPartitioner extends
    Partitioner<PairOfStrings, IntWritable> {
    @Override
    public int getPartition(PairOfStrings key, IntWritable value,
        int numReduceTasks) {
        return (key.getLeftElement().hashCode() & Integer.MAX_VALUE)
            % numReduceTasks;
    }
}
```

Partition based on the left element only

Reducer

Make sure $(a, *)$ comes first (already guaranteed)

Hold state in reducer across different key-value pairs

```
1: class REDUCER
2:   method REDUCE(pair  $p$ , counts  $[c_1, c_2, \dots]$ )
3:      $s \leftarrow 0$ 
4:     for all count  $c \in$  counts  $[c_1, c_2, \dots]$  do
5:        $s \leftarrow s + c$  ▷ Sum co-occurrence counts
6:     EMIT(pair  $p$ , count  $s$ )
```

Reducer

Make sure (a, *) comes first (already guaranteed)

Hold state in reducer across different key-value pairs

```
1: class REDUCER
2:   method REDUCE(pair  $p$ , counts  $[c_1, c_2, \dots]$ )
3:      $s \leftarrow 0$ 
4:     for all count  $c \in$  counts  $[c_1, c_2, \dots]$  do
5:        $s \leftarrow s + c$  ▷ Sum co-occurrence counts

  if  $p.\text{rightElement} == ""$  then
     $\text{marginal} \leftarrow s$ 
  else
    Emit(pair  $p$ , frequency  $s / \text{marginal}$ )
```

Where should *marginal* be declared and initialized?

You'll write the complete
code in Assignment-3

“Order inversion”

Common design pattern:

- ▶ take advantage of sorted key order at reducer to sequence computations
- ▶ get the marginal counts to arrive at the reducer before the joint counts

Optimization:

- ▶ apply in-memory combining pattern to accumulate marginal counts

Pairs vs. Stripes

Pairs

Turn synchronization into an ordering problem

- ▶ sort keys into correct order of computation
- ▶ partition key space so that each reducer gets the appropriate set of partial results
- ▶ hold state in reducer *across* multiple key-value pairs to perform computation

Stripes

Construct data structures that bring partial results together

- ▶ each reducer receives all data it needs to complete the computation
- ▶ be careful about the scalability issue: large stripes **overflow** the memory (and network)!
 - ▶ thus usually requires special treatment

Secondary sorting

Secondary sorting

In Hadoop, MapReduce sorts input to reducers by key

- ▶ values may be arbitrarily ordered
 - ▶ Google's proprietary implementation supports value sorting

What if we want to sort values as well?

- ▶ e.g., $k \rightarrow (v_1, r), (v_3, r), (v_4, r), (v_8, r) \dots$

Secondary sorting: solutions

Solution 1:

- ▶ buffer values in memory, then sort
- ▶ Is this a good idea? Why?

Solution 2:

- ▶ “value-to-key conversion” design pattern

“Value-to-key conversion”

Form composite intermediate key:

$$k \rightarrow (v_1, r), (v_3, r), (v_4, r), (v_8, r) \dots \Rightarrow \begin{array}{l} (k, v_1) \rightarrow r \\ (k, v_3) \rightarrow r \\ (k, v_4) \rightarrow r \\ (k, v_8) \rightarrow r \\ \dots \end{array}$$

Let execution framework do sorting

Preserve state *across* multiple key-value pairs to handle processing

Anything else we need to do?

Recap: tools for synchronization

Cleverly-constructed data structures

- ▶ bring data together

Sort order of intermediate keys

- ▶ control order in which reducers process keys

Partitioner

- ▶ control which reducer processes which keys

Preserving state in mappers and reducers

- ▶ capture dependencies *across* multiple keys and values

Issues and tradeoffs

Number of key-value pairs

- ▶ Object creation overhead
- ▶ Time for sorting and shuffling pairs across the network

Size of each key-value pair

- ▶ De/serialization overhead

Issues and tradeoffs

Local aggregation

- ▶ Opportunities to perform local aggregation varies
- ▶ Combiners make a big difference
- ▶ Combiners vs. in-mapper combining
- ▶ RAM vs. disk vs. network

Debugging at scale

Work on small datasets, won't scale... why?

- ▶ memory management issues (buffering and object creation)
- ▶ too much intermediate data
- ▶ mangled input records

Real-world data is messy!

- ▶ there's no such thing as "consistent data"
- ▶ watch out for corner cases
- ▶ isolate unexpected behavior, bring local

Credits

Slides are adapted from Prof. Jimmy Lin's slides at the University of Waterloo